Sports maker replay v0

Does the favourite–longshot table survive contact with phase, cost and queue?

Follow-on to the FLB stratified measurement. Same 15 sealed trading days, but the sports cells are now split pre-game vs in-game, priced as a one-lot maker who fades the taker, and charged the real worst-case bill.

+$5k	**−10¢**	**-$178**	**25–33**
fading favourites on props, 15 days, one lot per cell	worst single fill on the favourite line — structurally capped	the same trade on single-game longshots: the tax is gone	contracts of displayed depth on the prop favourite NO side — the real ceiling

Three pre-registered questions, three answers.

① Is favourite over-pricing a pre-game or an in-game phenomenon? Both, and the pre-game unit is fatter. Single-game favourites are over-priced by +11.4pp pre-game (n=182) against +2.8pp in-game (n≈1,300). That releases the FLB report's suspended verdict on this end: it is not just chasing a team that went ahead.

② What does the four-line worst-case bill look like? The two ends have opposite shapes. Fading favourites: small wins, high frequency, worst single fill capped at −10¢, drawdown ≈ 0. Fading longshots: positive on props, negative on single games, and every line carries a −94 to −99¢ tail. The +10~16pp lottery tax in the FLB table only lives inside the last ten minutes; aggregate the whole in-game phase and fair pricing earlier in the match dilutes it away.

③ Is there room to trade it? Not much, and the constraint is the queue, not the flow. The prop favourite NO side shows 25–33 contracts of depth against 49.7M contracts of volume in the same cells.

15 trading days (2026-07-10..24) · 346,410 replayed positions (135,386 in-game, 23,301 pre-game) · fill model = one lot inherited at the volume-weighted taker price, NOT a per-event queue simulation — Q3 measures how far that proxy sits from reality

Q1 · Is the favourite over-priced before the match, or only during it?

over-payment by phase

Taker over-payment = price paid − realized win rate, in percentage points. Positive means the taker overpaid and the maker on the other side collected it.

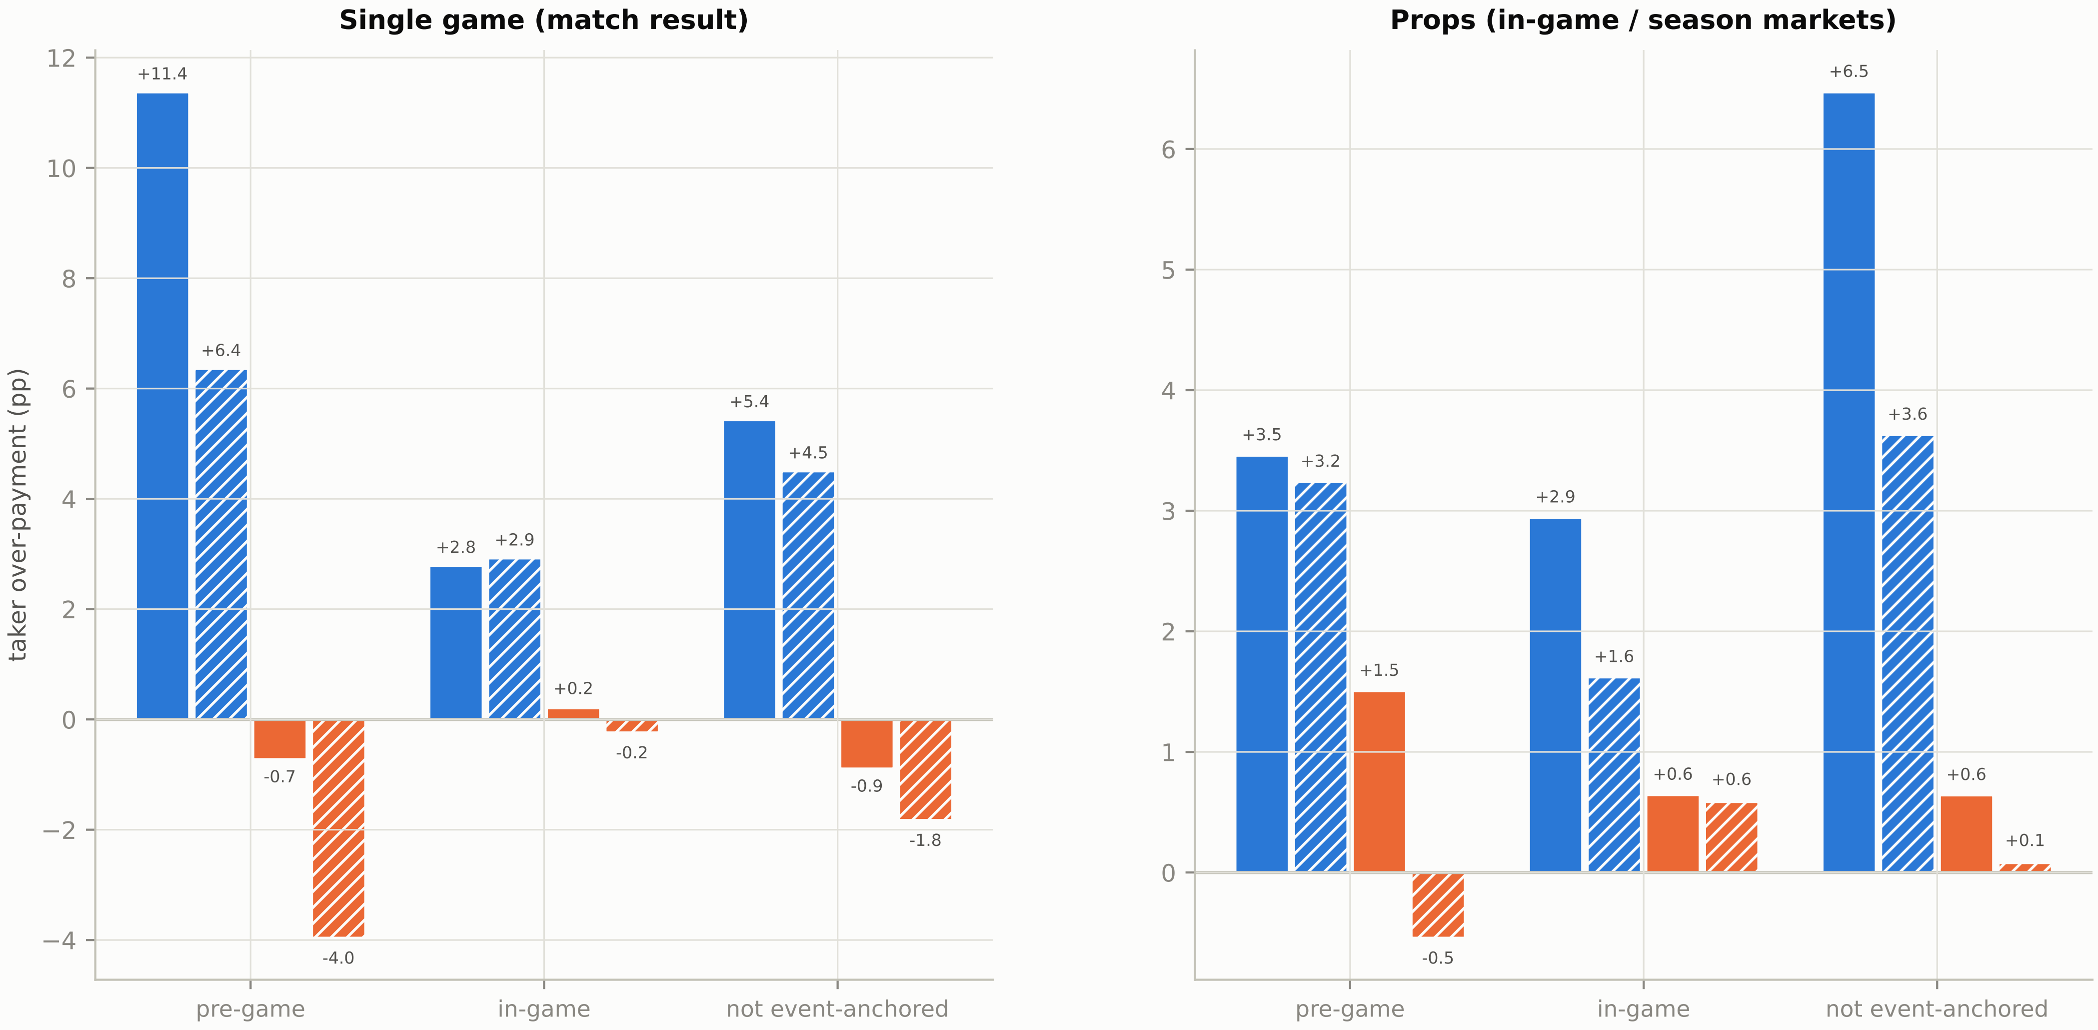

Hatched bars are the NO side of the same cell — a genuine second sample, not a mirror, because the two sides are traded by different flow. "Not event-anchored" is the residual: season-long and multi-event props with no single kick-off to split on. Read the favourite columns (blue): over-payment is present in every phase and every family, largest pre-game on single games (+11.4pp on the YES side, n=182). The longshot columns (orange) are near zero or negative once the whole phase is aggregated — which is the correction this replay exists to deliver.

Q2 · The four lines, day by day

cumulative P&L

Cumulative P&L of a one-lot maker fading each end, all phases aggregated. The shape matters more than the level: what we need is a line with no cliff in it.

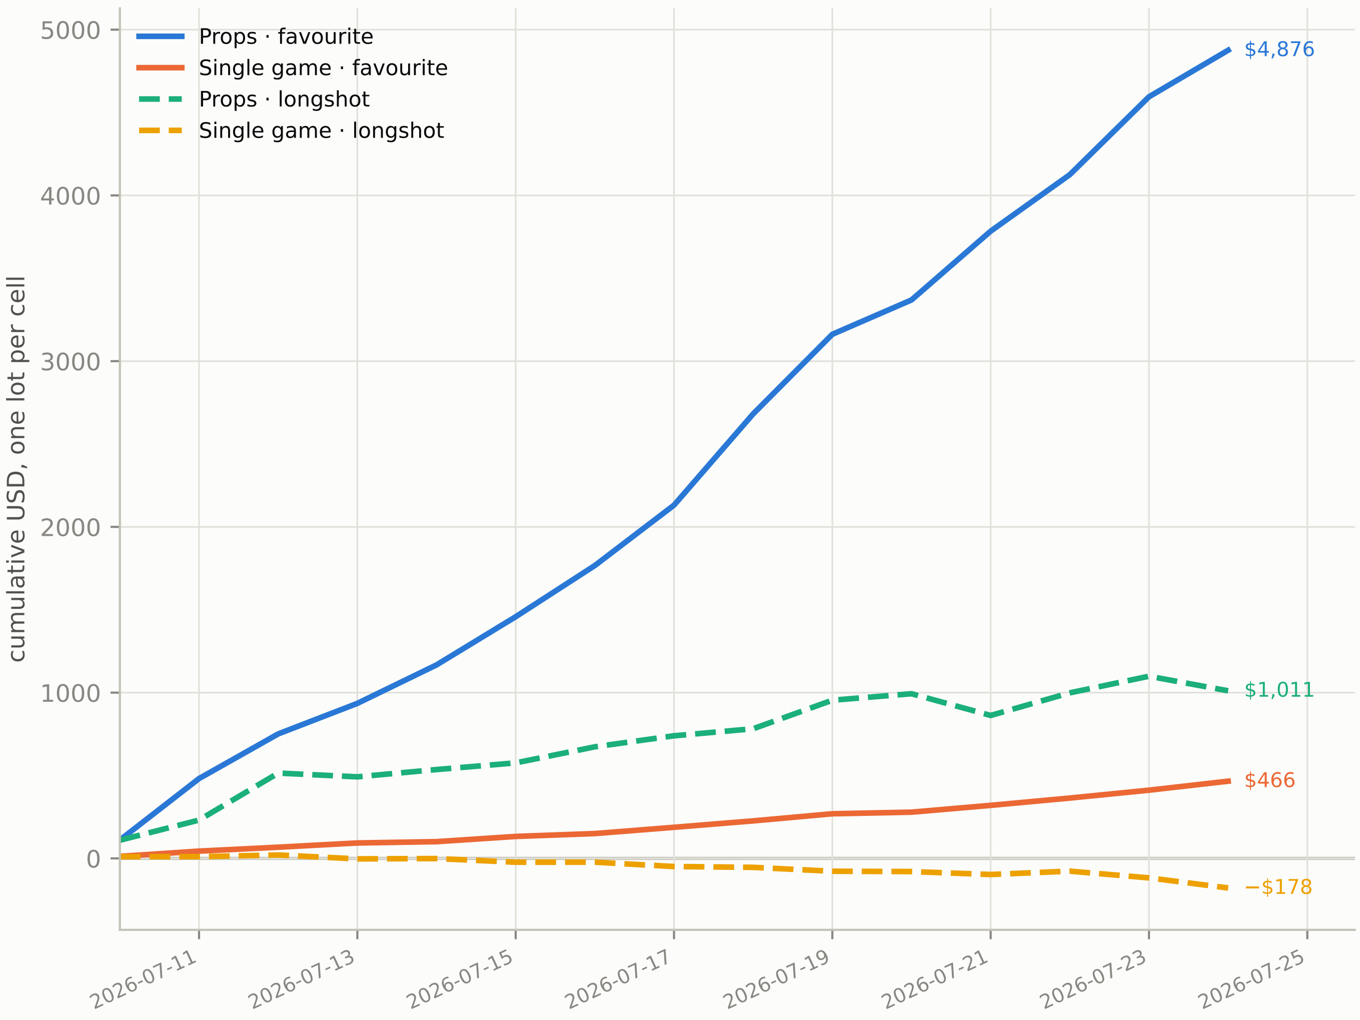

book	end	total	EV ¢/lot	worst fill	worst day	max draw down	fills
Props	favourite	$4,876	+3.59	-10¢	−$1	−$0	135,565
Single game	favourite	$466	+4.12	-10¢	−$4	−$4	11,294
Props	longshot	$1,011	+0.57	-99¢	−$54	−$75	177,017
Single game	longshot	−$178	-0.79	-99¢	−$37	−$149	22,534

One lot per (market, cell) — no sizing, no compounding, no inventory. The two favourite lines climb in small steps and never give it back; the two longshot lines are step functions with cliffs, which is the same −94 to −99¢ tail showing up as realized days. The single-game longshot line ends below where it started. Loss-budget doctrine reads the last three columns, not the first: a line whose worst single fill is structurally capped at −10¢ is a different animal from one that can print −99¢.

Q2b · The worst-case bill, per line and phase

loss budget

The loss-budget doctrine wants four rows filled in before anything goes live: worst single fill, worst day, max drawdown, and the line total. Here they are.

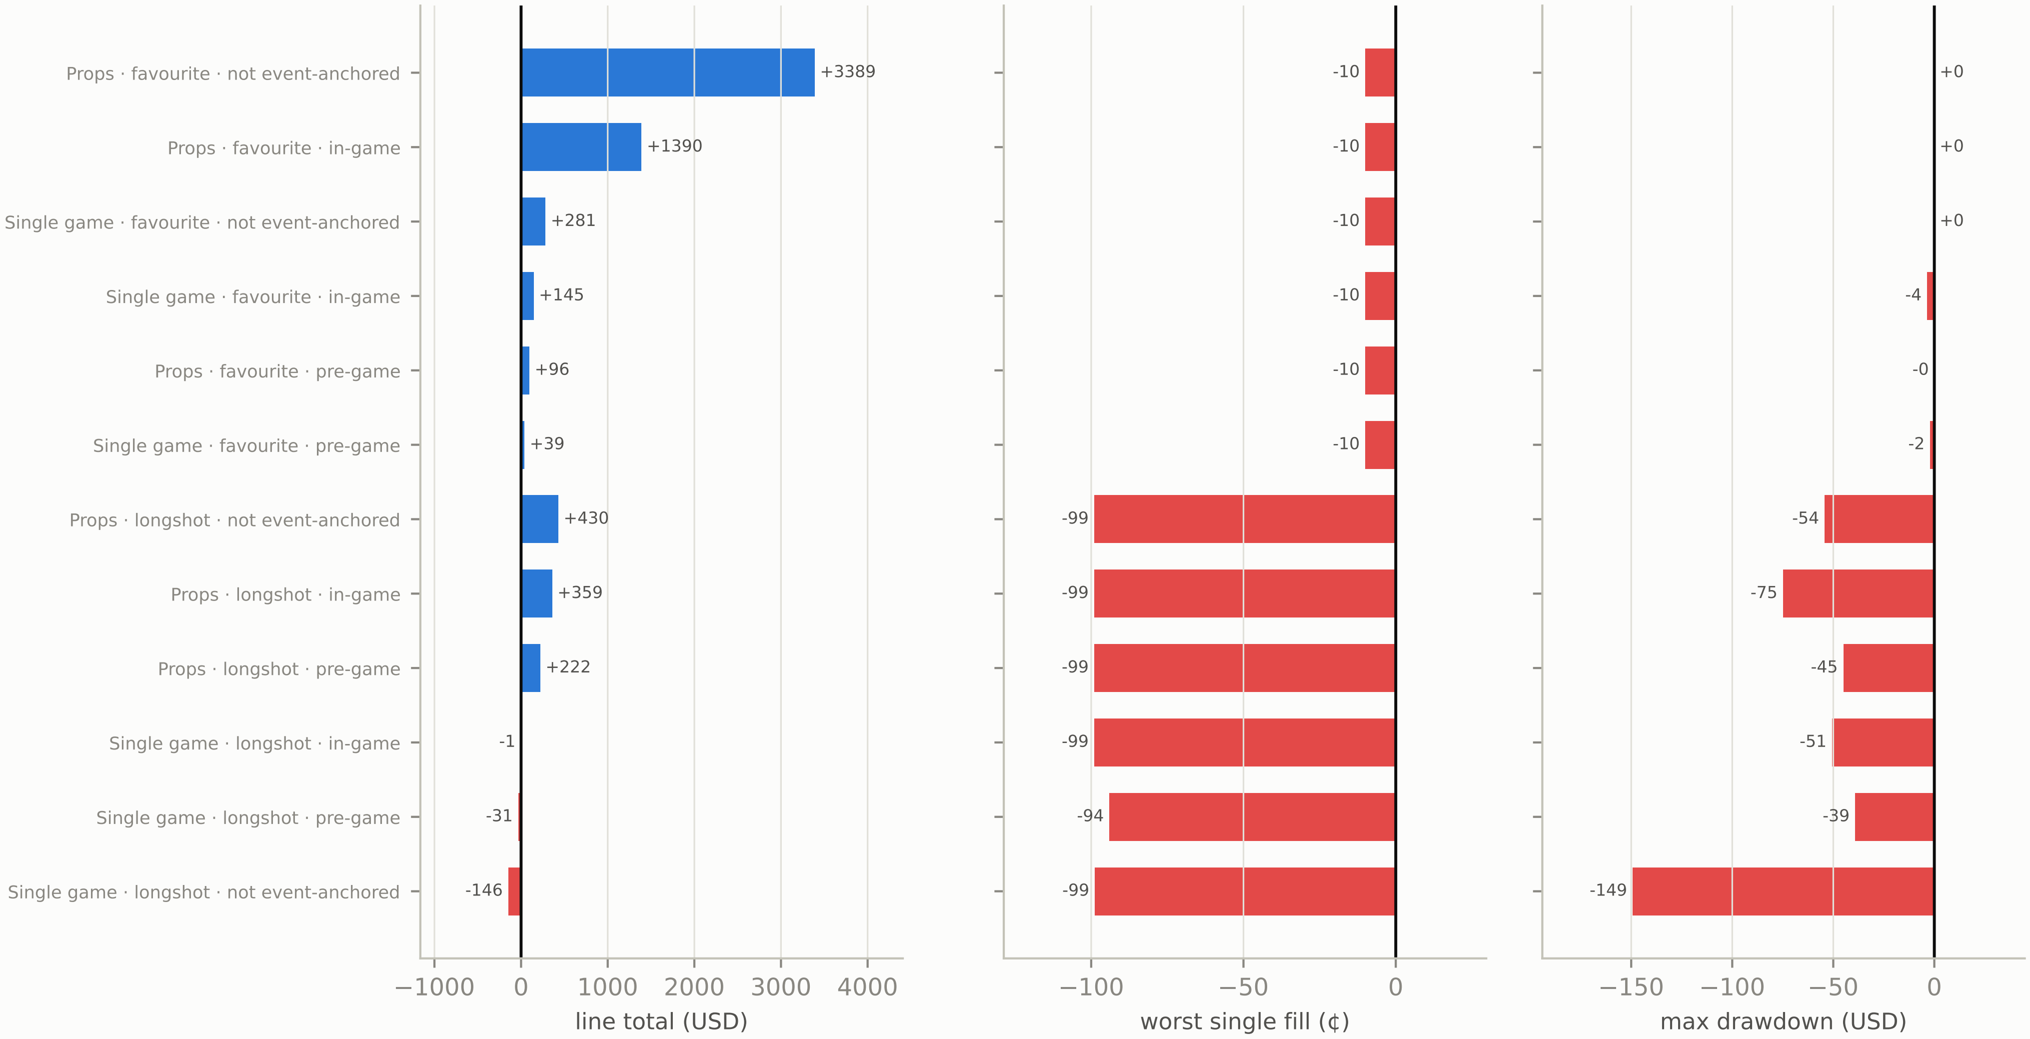

Sorted with the favourite lines on top. The favourite block is uniform: every row is positive, every worst single fill is exactly −10¢ (the fade is capped by how far a 90–99¢ contract can move against the maker), and drawdown is zero or a rounding error. The longshot block has the opposite signature — a −94 to −99¢ worst fill in every row, drawdowns to −$149, and one line negative outright. Same table, two completely different risk objects; only the first survives the doctrine.

Q3 · How much of it could we actually get?

queue reality

Median displayed depth per price bucket, by book and side. This is the ceiling on the replay above — the one-lot model quietly assumes a fill that the queue has to provide.

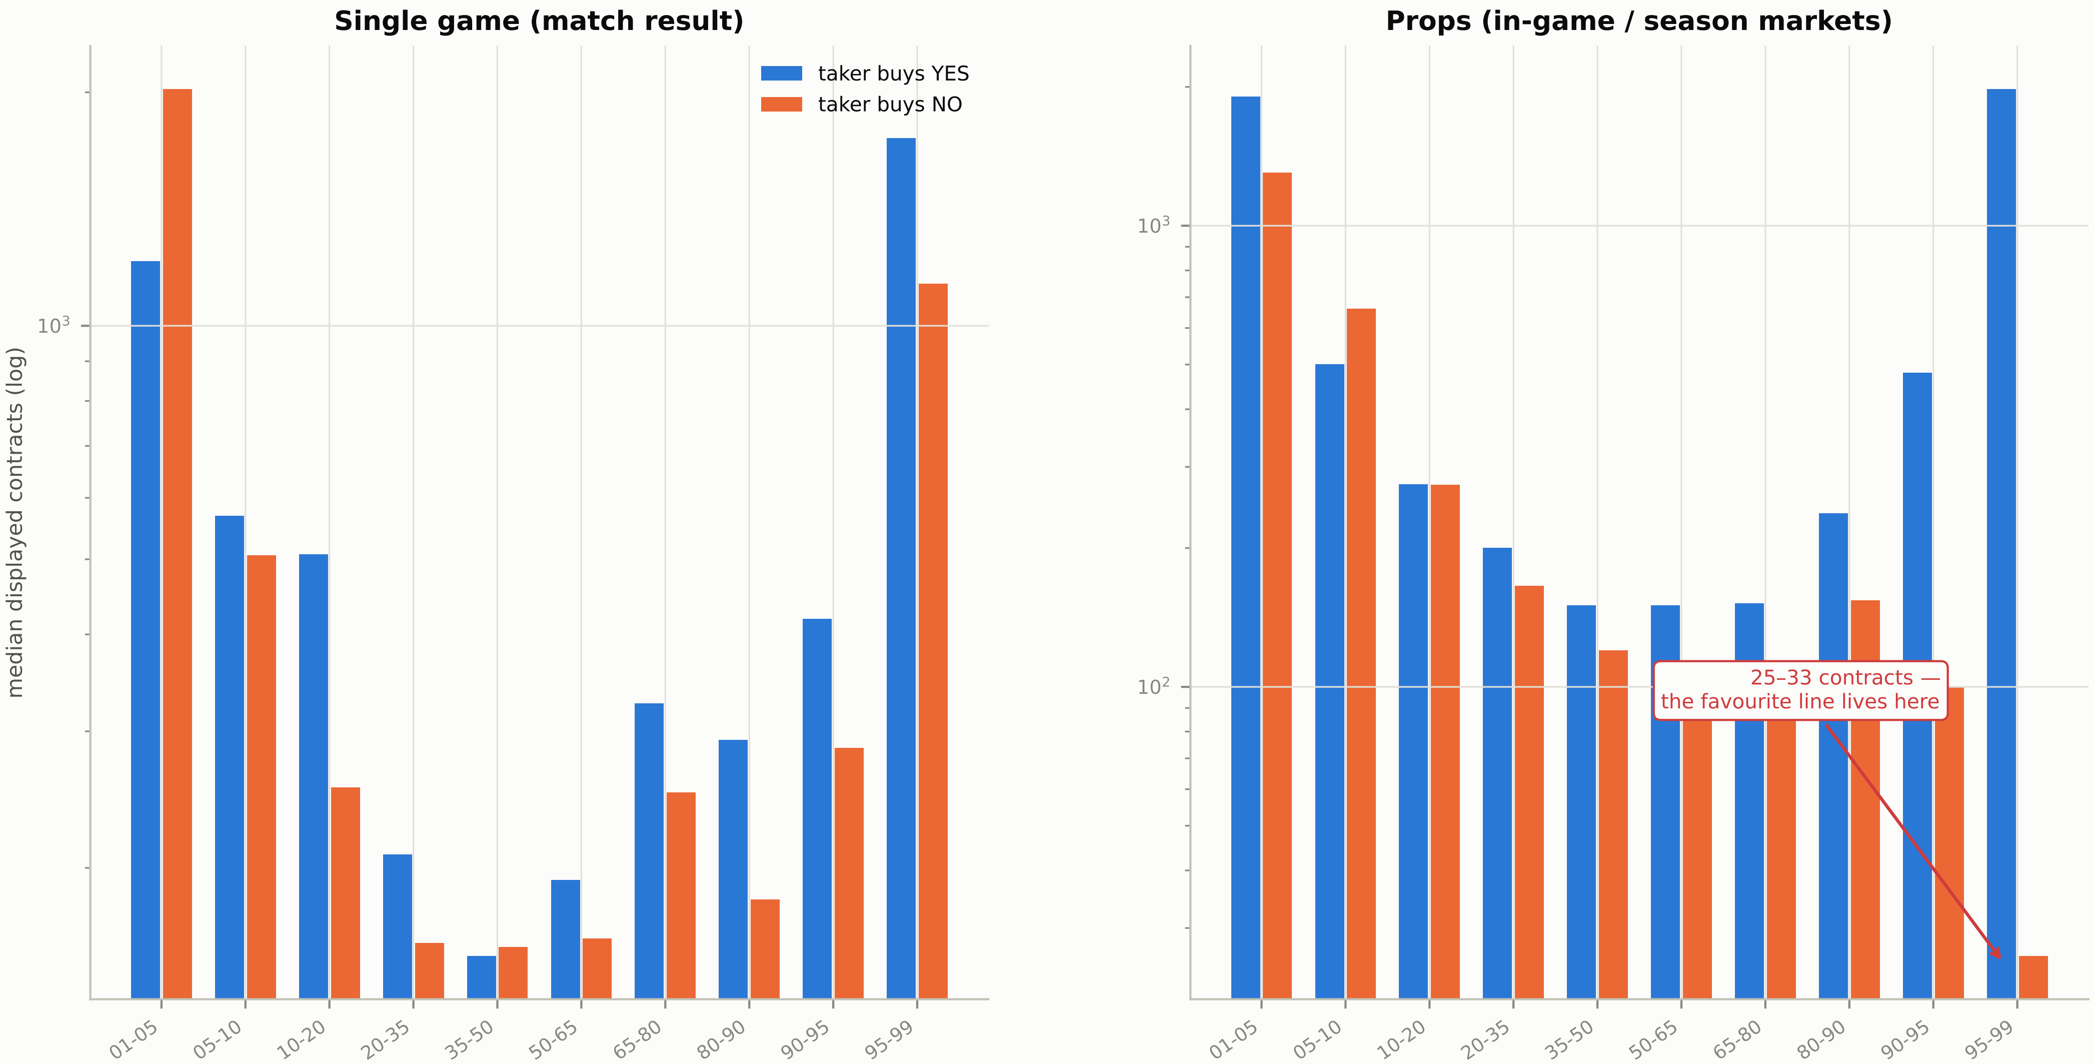

Log scale — the spread across buckets is two orders of magnitude. The number that matters is the prop NO side at 90–99¢: a median of 25–33 displayed contracts, against 49.7M contracts of volume traded in the same cells over the window. Capacity is queue-bound, not flow-bound. Quoting 10 lots there takes roughly a quarter of the queue pro-rata; the YES side (767 in-game / 2,754 pre-game) dilutes our share to a few per cent. Any size extrapolation from the one-lot replay has to pass through this chart first.

Method, deviations, and what this is not

What was replayed

Every sealed-day sports print that joined settlement truth was assigned to a price bucket and a phase, and a one-lot maker position was booked on the opposite side of the taker at the cell's volume-weighted taker price, held to settlement. 346,410 positions across 164,453 markets. Phase comes from the event anchor: pre-game before kick-off, in-game after it, and "not event-anchored" for season-long and multi-event props that have no single kick-off — that residual is the largest block (187,723 positions) and is reported separately rather than folded into either.

The fill model, stated plainly

This is a one-lot proxy that inherits the market's volume-weighted taker price. It is NOT a per-event queue simulation. It assumes we were at the front of the queue for one contract every time a taker crossed, which is optimistic in exactly the cells where the edge is largest. Q3 exists to measure how far that assumption sits from the displayed book. Fees are not netted in this replay— see the Who Profits report for the maker-fee multipliers, which bite on precisely these per-game sports series.

What this corrects in the FLB report

The stratified table showed a +10~16pp "lottery tax" on sports longshots inside the last ten minutes and suspended judgement on the time dimension because the convergence check failed. Splitting by phase resolves both. The favourite end is real, pre-game as well as in-game, so its verdict is released. The longshot end is a narrow-band artifact: aggregate the whole in-game phase and it nearly vanishes, and on single games it goes negative. The convergence failure was end-of-game behaviour, not a broken anchor— settlement_ts is never earlier than close_time anywhere in the warehouse.

What it would take to trade this

Four things, none of them done here: a per-event L2 queue simulation (we have the full book, so this is buildable); robustness across a second window, since 15 days inside one season is not a prior; the worst-case bill recomputed at real size and submitted for approval under the loss-budget doctrine; and registration as an RC line. The arithmetic extrapolation in the source report (+$6k/day gross at 20 lots) is arithmetic only — it ignores that larger size concentrates fills at the moments the queue gets swept, which is adverse selection by another name.

Source report: agents/audit/REPORT_SPORTS_MAKER_REPLAY_2026-07-27.md · code: tools/research/replication/experiments/sports_maker_replay.py · artifacts: Repo Research/FLB_STRATIFIED_20260727/sports_maker_replay/