

Who Profits on Kalshi?

Execution, not forecasting — a replication of Della Vedova (2026) on our own exchange tape

Della Vedova (2026) says the price you pay decides who profits, not the side you pick. We ran his four tests on our own sealed Kalshi tape — 17 days, the whole exchange, six market families kept strictly separate.

−\$235M

what liquidity takers lost, net of fees, in 15 settled days across the six families

+\$113M

what liquidity makers made over the same prints, net of maker fees

\$122M

what Kalshi took in fees — more than the entire maker profit pool it sits beside

12.8¢ → 1.0¢

what ten points of extra forecasting accuracy is worth at 50¢ versus at 95¢


What replicates, and what does not.

- ① The paper's headline mechanism replicates almost exactly. A ten-point gain in directional accuracy is worth 11.8-13.5¢ per contract when price sits at 50¢ and 0.5-1.8¢ once price reaches 95¢ (paper: 11.6¢ → 1.5¢). The interaction coefficient is −2.31 to −2.89 across our six families against the paper's −2.244, on a different exchange, a different asset mix and a different year. Being right stops paying once the price has already moved; arriving early is what pays.
- ② The benchmark-contamination result replicates cleanly. Regress execution edge on directional edge and the coefficient is −0.10 to −0.34 under a backward-VWAP reference price, and −0.000 to −0.037 once the reference is a single adjacent print (paper: −0.178 → 0.000). The apparent trade-off between picking well and executing well is an artifact of scoring flow against a benchmark built from that same flow.
- ③ The inversion does not replicate in the paper's form, and the reason is instructive. Polymarket retail picks winners 51.3% of the time and still loses. Our takers pick the winning side only 42-49% of the time (contract-weighted) — they lose on the side picked as well as on the price paid, because Kalshi taker flow is tilted hard into cheap longshots. Same conclusion, blunter: the profitable role is not the accurate one.
- ④ Kalshi has a fee, and Polymarket did not. That single difference dominates our tape: takers pay 0.6-1.1¢ per contract, which is larger than the entire gross edge the makers on the other side of those same prints collect. Kalshi's fee take, \$122M in 15 days, exceeds every dollar the maker side kept.

E1 · The inversion

accuracy vs return

Paper Figure 1 analog. Each point is one market family × order role. The role that picks winners more often is the role that loses money.




Marker area $\propto$ contracts traded. Both roles are net of the fees that role actually pays: takers $0.07 \cdot p \cdot (1-p)$ ¢/contract on every series, makers $0.0175 \cdot p \cdot (1-p)$ on the ~80 series Kalshi lists with a maker multiplier and nothing elsewhere. Capital at risk = $\sum$ price paid $\times$ size on that role's own side. Zero-sum before fees: maker gross P&L = -taker gross P&L by construction, so the gap between the two bars is exactly the exchange's cut.

E2 · Nobody beats the price

excess accuracy

Paper Table 2 Panel B analog. Excess accuracy = realized win rate – max(p, 1–p), the win rate a no-skill trader gets by always siding with the favourite.

Both sides of every print sit below the price-implied benchmark




family	raw accuracy	price-implied	excess (taker)	mean price paid
Crypto 15-min	48.2%	75.0%	-26.8 pp	48.5¢
Crypto hourly/daily	43.6%	78.3%	-34.6 pp	44.1¢
Sports game	48.6%	75.7%	-27.1 pp	49.0¢
Sports props	42.1%	77.5%	-35.3 pp	43.8¢
Politics/elections	27.6%	87.6%	-60.0 pp	28.4¢
Economics	40.1%	87.5%	-47.4 pp	43.2¢

Raw accuracy above 50% is not skill: a contract bought at 70¢ wins ~70% of the time with no information at all. Excess accuracy nets that out. Because max(p,1–p) is symmetric in the two sides of a print, maker and taker excess accuracy sum to 1–2·max(p,1–p) ≤ 0 — mechanically, at most one side can beat the price, and here neither does.

E3 · Where the money actually comes from

edge decomposition

Paper Figure 2 analog. Taker edge per contract split into a directional part (outcome – fair) and an execution part (fair – price paid), fair = backward VWAP.




Backward VWAP is the paper's primary benchmark and is contaminated by the evaluated flow's own prior prints— see E4 for what that does to the split between the two columns. The bottom line (net ¢/contract) does not depend on the benchmark: it is outcome – price paid – fee, measured against settlement.

E4 · Benchmark contamination

cross-skill coefficient

Paper Table 4 analog. The same regression — execution edge on directional edge, one observation per market × role — run under three reference prices that draw progressively less on the evaluated flow's own prints.









Paper: $\beta = -0.178$ under backward VWAP $\rightarrow -0.013$ under local VWAP $\rightarrow 0.000$ under the lagged price, and the analytical bias bound reproduces the contaminated number to within 1%. A real relation between the two skills would not move when only the reference price changes. Caveat: our unit is market × order role, not wallet — Kalshi's tape carries no trader identity — so this measures the mechanical benchmark bias, which is exactly what the paper attributes the negative coefficient to; it is not a wallet-level skill correlation.

E5 · Lifecycle timing — when the edge is there

edge at entry

Paper Figure 5 analog. Edge per contract by how long before the market closed the print happened. Both series are net of the fee that role pays.




Time bands are minutes before market close (settlement anchor = close_time, fallback settlement_ts). The paper finds bot edge-at-entry decays monotonically over the market lifecycle — large early, ~zero at 7-14 days, negative at the close — because the value of being right collapses as price converges on the outcome (E6). Read this chart as where in a market's life liquidity provision is paid, and where it is picked off.

E6 · Why execution wins: being right stops paying

accuracy × price distance

Paper Table 7 analog. Edge regressed on accuracy, on distance of the traded price from 50¢, and on their interaction (market × role observations).




family	β accuracy	β interaction	R²
Crypto 15-min	1.352	-2.892	0.94
Crypto hourly/daily	1.283	-2.604	0.90
Sports game	1.255	-2.514	0.91
Sports props	1.182	-2.305	0.88
Politics/elections	1.251	-2.482	0.86
Economics	1.355	-2.619	0.67
paper (Polymarket)	1.162	-2.244	0.78

Coefficients in dollars per contract. A negative interaction means a correct forecast is worth progressively less the further price has already moved toward the outcome — the binary payoff ceiling. In the paper a ten-point accuracy gain is worth 11.6¢ at 50¢ and 1.5¢ at 95¢, an eightfold collapse; that is why the trader who arrives early with no view beats the trader who arrives late with a good one.

E7 · Are Kalshi prices honest probabilities?

miscalibration by family

Realized win rate minus the price paid, per price bucket. Zero means the price was a fair probability. The dashed line is the taker fee — anything below it loses money even when the price is fair.




One point per price bucket, plotted as the deviation from a fair price rather than on a 0-100 diagonal, where deviations this small are invisible. Error bars are 95% cluster bootstrap over markets (1,000 resamples, cluster = market— the house rule, because per-print n inflates by tens of times). Two things to read off it. First, the tilt is the textbook favourite-longshot bias: cheap contracts settle less often than their price promises (points below zero on the left) and expensive ones more often (points above zero on the right). Second, and more important for us, the fee curve sits above almost every point— Kalshi prices are close enough to honest that lifting a fair contract is a losing trade, and selling into one is a winning trade, before either side has an opinion.

E8 · Adverse selection: how fast the maker's edge decays


markout vs settlement

Maker edge measured two ways on the same fills: mark-to-market 5 minutes after the fill, and carried to settlement.

Maker edge at two horizons



markout over the lifecycle




Markout = the market's own price 0-5 minutes after the fill, evaluated on the side the maker took, minus the maker's entry price. A positive markout is spread captured and kept; the gap between markout and settlement is what the position gives back (or gains) over the rest of the market's life. Coverage: fills with at least one later print inside the 5-minute window (Crypto 15-min 100%, Crypto hourly/daily 98%, Sports game 99%, Sports props 97%, Politics/elections 83%, Economics 56%).

E9 · Which books actually pay a market maker

maker edge net of fee

The 18 highest-volume series on the tape. Blue is the maker's gross edge per contract; the hatched part is the series' own maker fee, which Kalshi charges on roughly eighty listed series and on no others.

maker P&L



Sorted by net edge. A wide confidence interval means the series' maker profit came from a handful of markets, not from a repeatable edge— treat those as untraded. The tennis-match books (KXATPMATCH, KXWTAMATCH) are the clearest warning: gross maker edge is real but the $0.0175 \cdot p \cdot (1-p)$ maker fee is the same size, so the net is zero or below. The fee-free crypto books are the opposite case.

Numbers and what they mean for us

summary

family	markets	prints	contracts	taker accuracy	excess accuracy	taker gross ¢/ct	taker fee ¢/ct	taker net ¢/ct	maker gross ¢/ct	maker fee ¢/ct	maker net ¢/ct [95% CI]	taker P&L	maker P&L	fees to exchange
Crypto 15-min	2,741	32,758,479	2,014,960,656	48.2%	-26.8 pp	-0.27	1.14	-1.41	+0.27	0.00	+0.27 [0.11, 0.44]	-\$28.49M	\$5.50M	\$22.99M
Crypto hourly/daily	25,572	6,977,890	686,728,755	43.6%	-34.6 pp	-0.47	1.03	-1.50	+0.47	0.00	+0.47 [0.25, 0.68]	-\$10.29M	\$3.25M	\$7.04M
Sports game	5,100	16,244,399	2,274,001,270	48.6%	-27.1 pp	-0.39	1.11	-1.49	+0.39	0.21	+0.18 [-1.17, 1.57]	-\$33.99M	\$4.07M	\$29.92M
Sports props	87,507	41,068,634	5,850,513,721	42.1%	-35.3 pp	-1.72	1.03	-2.76	+1.72	0.03	+1.70 [-0.85, 4.35]	-\$161.22M	\$99.20M	\$62.02M
Politics/elections	3,288	1,334,407	82,711,685	27.6%	-60.0 pp	-0.79	0.61	-1.40	+0.79	0.00	+0.79 [-0.61, 1.86]	-\$1.15M	\$652k	\$502k
Economics	807	132,019	5,593,557	40.1%	-47.4 pp	-3.08	0.63	-3.71	+3.08	0.02	+3.06 [1.00, 5.73]	-\$208k	\$171k	\$36k

For our own book — what these numbers say we should do.

- Quote where the maker edge survives its own fee. Crypto 15-minute (+0.27¢/contract, 95% CI [0.11, 0.44]) and crypto hourly/daily (+0.47¢, CI [0.25, 0.68]) are the two books where the maker edge is both positive and tight across markets — and Kalshi charges no maker fee there. That is a per-market-repeatable edge, not a handful of lucky markets.
- Read the maker-fee list before quoting a sports book. Kalshi charges makers $0.0175 \cdot p \cdot (1-p)$ on ~80 listed series and nothing elsewhere. On KXATPMATCH the maker fee (0.277¢) is the whole gross edge (0.271¢) — net zero. On KXWTAMATCH the net is -0.34¢. Sports game maker edge is +0.39¢ gross but +0.18¢ after fee, with a confidence interval straddling zero.
- Do not read the big sports-prop maker profit as an edge. It is +\$99M, but 78% of it sits in five series (World Cup outright, UFC fights, World Cup goal markets) and the per-market interval is [-0.85, +4.35]¢ — that is a favourite-longshot windfall in a few crowded books, not a strategy. The paper hits the same wall: its top-10 bot wallets earn 70% of all bot profit.
- Stop quoting into the close. Maker edge decays monotonically toward expiry and turns negative in the final minutes — crypto 15-min goes +0.66¢ (10-60 min out) → +0.26¢ (5-10 min) → +0.01¢ (2-5 min) → -0.35¢ (under 2 min). The last band is where our fills are worth least and our adverse selection is worst. Whatever we quote there needs a materially wider spread or no quote at all.
- Stop scoring our own execution against a VWAP we helped set. Every execution-quality number we produce from a backward VWAP built out of our own prints carries the bias measured on page 5. Use the adjacent print, or a forward window that skips our own fills.

Caveats, stated plainly. ► Kalshi's public tape carries no trader identity, so the paper's wallet-level classification (bot / sophisticated / retail / casual) and its within-wallet persistence tests cannot be run at all; our axis is the order role recorded on each print, which is the paper's own mechanism margin, not its classification margin ► With two roles per print the maker's components are the exact mirror of the taker's, so "maker vs taker" here is one comparison, not two independent populations ► Settlement truth comes from a catalog snapshot taken 2026-07-24T21:30Z: crypto and sports join at 98-100%, politics at 79% and economics at 60%, so those two families are selected toward markets that resolved early — read them as indicative. ► For the four high-volume families the backward-VWAP benchmark is rebuilt each day rather than across a market's whole life; their markets are intraday, so this costs each market its opening print. ► Fees are modelled continuously from the published July 7 2026 schedule (taker $0.07 \cdot p \cdot (1-p)$, maker $0.0175 \cdot p \cdot (1-p)$ on listed series only); Kalshi rounds up per order, so the true taker cost is a shade higher than shown. ► Markout uses subsequent trade prints, not quotes.